

# Gender Prediction

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [5]: data=pd.read_csv("C:/Users/rockz/OneDrive/Documents/Gender Prediction.csv")
```

```
In [6]: data
```

```
Out[6]:
```

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	dista
0	1	11.8	6.1	1	0	1	
1	0	14.0	5.4	0	0	1	
2	0	11.8	6.3	1	1	1	
3	0	14.4	6.1	0	1	1	
4	1	13.5	5.9	0	0	0	
...	...	...	...	...	...	...	...
4996	1	13.6	5.1	0	0	0	
4997	1	11.9	5.4	0	0	0	
4998	1	12.9	5.7	0	0	0	
4999	1	13.2	6.2	0	0	0	
5000	1	15.4	5.4	1	1	1	

5001 rows × 8 columns



```
In [7]: data.columns
```

```
Out[7]: Index(['long_hair', 'forehead_width_cm', 'forehead_height_cm', 'nose_wide',
'nose_long', 'lips_thin', 'distance_nose_to_lip_long', 'gender'],
dtype='object')
```

```
In [8]: data.describe()
```

```
Out[8]:
```


	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_
count	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000	5001.000000
mean	0.869626	13.181484	5.946311	0.493901	0.507898	0.493901
std	0.336748	1.107128	0.541268	0.500013	0.499988	0.500013
min	0.000000	11.400000	5.100000	0.000000	0.000000	0.000000
25%	1.000000	12.200000	5.500000	0.000000	0.000000	0.000000
50%	1.000000	13.100000	5.900000	0.000000	1.000000	0.000000
75%	1.000000	14.000000	6.400000	1.000000	1.000000	1.000000
max	1.000000	15.500000	7.100000	1.000000	1.000000	1.000000

```
In [9]: data['gender'].value_counts()
```

```
Out[9]: Female    2501
Male          2500
Name: gender, dtype: int64
```

```
In [57]: plt.figure(figsize=(8,6))
sns.countplot(x="gender",data=data,palette="twilight")
```

```
Out[57]: <AxesSubplot:xlabel='gender', ylabel='count'>
```



```
In [11]: plt.figure(figsize=(15,10))
sns.heatmap(data.corr(),annot=True,linewidths=0.5,cmap="Reds")
```

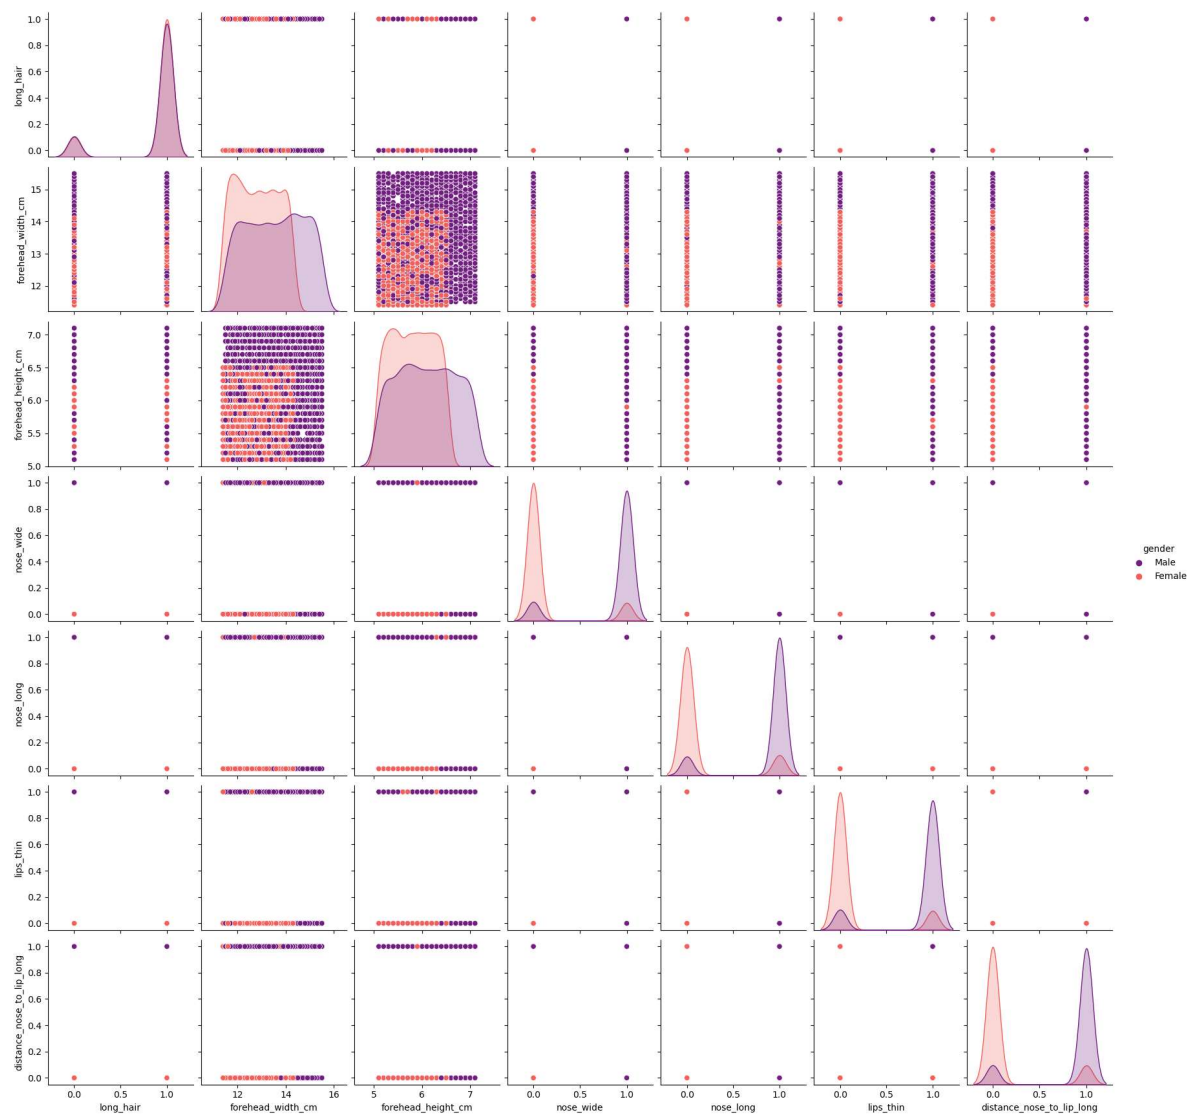
Out[11]: <AxesSubplot:>



```
In [12]: m_col=['long_hair','forehead_width_cm','forehead_height_cm','nose_wide','nose_
```

```
In [13]: sns.pairplot(data[m_col],hue='gender',palette='magma')
```

```
Out[13]: <seaborn.axisgrid.PairGrid at 0x25a1285b1f0>
```



```
In [14]: x=data.drop('gender',axis=1)
y=data['gender']
```

In [15]: x

Out[15]:

	long_hair	forehead_width_cm	forehead_height_cm	nose_wide	nose_long	lips_thin	dista
0	1	11.8	6.1	1	0	1	
1	0	14.0	5.4	0	0	1	
2	0	11.8	6.3	1	1	1	
3	0	14.4	6.1	0	1	1	
4	1	13.5	5.9	0	0	0	
...	...	...	...	...	...	...	...
4996	1	13.6	5.1	0	0	0	
4997	1	11.9	5.4	0	0	0	
4998	1	12.9	5.7	0	0	0	
4999	1	13.2	6.2	0	0	0	
5000	1	15.4	5.4	1	1	1	

5001 rows × 7 columns



In [16]: y

Out[16]:

0	Male
1	Female
2	Male
3	Male
4	Female
...	...
4996	Female
4997	Female
4998	Female
4999	Female
5000	Male

Name: gender, Length: 5001, dtype: object

```
In [25]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.40,random_state=42)
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix,classification_report
```

## LOGISTIC REGRESSION

```
In [21]: from sklearn.linear_model import LogisticRegression
log=LogisticRegression()
log.fit(x_train,y_train)
```

Out[21]: LogisticRegression()

```
In [33]: logpredict=log.predict(x_train)
log_acc=accuracy_score(y_train,logpredict)
log_acc
```

Out[33]: 0.9673333333333334

```
In [26]: confusion_matrix(y_train,logpredict)
```

Out[26]: array([[1453, 47],  
[ 51, 1449]], dtype=int64)

```
In [28]: classification_report(y_train,logpredict)
```

Out[28]:

	precision	recall	f1-score	support	Female
0.97	0.97	0.97	1500	Male	0.97
97	1500	accuracy	0.97	3000	mac
ro avg	0.97	0.97	0.97	3000	weighted avg
0.97	0.97	3000			

## K NEAREST NEIGHBORS

```
In [31]: from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier()
knn.fit(x_train,y_train)
```

Out[31]: KNeighborsClassifier()

```
In [34]: knnpredict=knn.predict(x_train)
knn_acc=accuracy_score(y_train,knnpredict)
knn_acc
```

C:\Users\rockz\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

Out[34]: 0.976

```
In [35]: confusion_matrix(y_train,knnpredict)
```

Out[35]: array([[1476, 24],  
[ 48, 1452]], dtype=int64)








```
In [55]: plt.figure(figsize=(8,6))  
model_acc=[log_acc,knn_acc,ran_acc,svm_acc]  
model_name=["LogisticRegression","KNN","RandomFoest","SVM"]  
sns.barplot(x=model_acc,y=model_name,palette="Oranges")
```

Out[55]: <AxesSubplot:>



## CONCLUSION

Random Forest Model gave best performance with an accuracy of 99.86%